



Inspector Mandi



Attempter Instructions

Welcome to the Inspector Mandi Project!

Welcome to the Inspector Mandi project! In this project, you will **evaluate AI-generated analytical reports** by comparing two HTML reports side-by-side and scoring them on two key dimensions: **visual quality** and **narrative excellence**.

Contents

 [Contents](#)

 [Instructions Change Log](#)

 [Project Overview:](#)

 [What You'll Do](#)

 [Step I: Reviewing the Task & Generating the UI](#)

[A. Generating & Viewing the Reports](#)

[B. Cross-Checking with an External Browser](#)

 [Step II: Visual Scoring](#)

[A. The Scoring Scale](#)

[B. The Standard to Keep in Mind](#)

[C. Calibration Rules](#)

[D. Issue Severity Reference](#)

- [E. A Note on Leniency](#)
-  [Step III: Visual Scoring](#)
- [A. Score Each Report Visually \(Repeat for Report A and Report B\)](#)
 - [B. Write the Visual Side-by-Side Comparison](#)
 - [C. Pillar 1 Scoring Criteria](#)
 - [D. Visual Scoring Guidelines](#)
-  [Step IV: Narrative Scoring](#)
- [A. Score Each Report on Narrative \(Repeat for Report A and Report B\)](#)
 - [B. Write the Narrative Side-by-Side Comparison](#)
 - [C. Pillar 2 Scoring Criteria](#)
 - [D. Narrative Scoring Guidelines](#)
-  [Step V: Overall Preference & Final Submission](#)
-  [Step VI: Common Errors](#)
- [Error 1: Leniency in Model Scoring](#)
 - [Error 2: Misclassifying Text Walls as a Visual Issue](#)
 - [Error 3: Writing Generic Rationales](#)
 - [Error 4: Scoring ASCII Tables Incorrectly](#)
 - [Error 5: Uncentered Tables with a Loose Left Column](#)
 - [Error 7: Markdown Rendering Failures](#)
 - [Error 8: Suboptimal Hyperlink Integration](#)



Instructions Change Log

Change Log		
Date	Summary Of Changes	Instructions Subject
May 27, 2026	First release	<i>Attempter Instructions</i>
May 29, 2026	Added Common Errors and updated the text wall from visual to narrative only	<i>Common Errors</i>
Jun 2, 2026	Updated rules (Don't penalize LaTeX rendering issue and mermaid charts). Added 2 sections in common errors.	<i>Visual and Common Error</i>
Jun 3, 2026	Added Step II, to explain scoring rules across the project with more detail. Update rule to avoid including comparisons between reports in the individual report justifications	<i>Understanding the Scoring Scale</i> <i>Visual, Narrative sections</i> <i>Common Errors</i>

Project Overview:

We would like to evaluate the outputs produced by an agent when tasked with **Deep Analysis and Reporting tasks**. These tasks require an agent to perform extensive research on any given topic using (optional) user-provided files and/or web search, carefully synthesize complex information from across the web and the files, and generate a **comprehensive, insightful text-based report** with appropriate visualizations. The target audience is professional. Reports should meet the standard of a polished business or analytical deliverable: management consulting, data science, or executive briefings. The expected output of the agent is always a report. 💡

What You'll Do

Step	Task	What to Complete
 Step I Review Task	Read the prompt and view both HTML reports	Read the prompt carefully Generate UI for each report Note any rendering discrepancies
 Step II: Understanding the Scoring Scale	Understand the scoring scale we are using in this project	Internalize the scoring rules.
 Step II Visual Scoring	Score Pillar 1 for Report A and Report B	Visual score (0–10) + rationale for Report A Visual score (0–10) + rationale for Report B Write visual side-by-side comparison
 Step III Narrative Scoring	Score Pillar 2 for Report A and Report B	Narrative score (0–10) + rationale for Report A Narrative score (0–10) + rationale for Report B Write narrative side-by-side comparison
★ Step IV Final Preference	Submit overall preference and summary	Select A/B preference on 7-point scale  Overall Preference


CONFIDENTIAL muzhaqu@gmail.com+outlier 2026-06-04T05:00:27.907Z	CONFIDENTIAL muzhaqu@gmail.com+outlier 2026-06-04T05:00:27.907Z	Rationale  Summary (1–2 sentences)  Feedback
---	---	--

Step I: Reviewing the Task & Generating the UI

In this step, your goal is to open and carefully inspect both reports before scoring. The reports are already pre-rendered — you do not need to generate anything. You are working directly with **Report A** and **Report B**. Take your time to view each report in full and note any visual defects, structural issues, or rendering problems before moving on to scoring.

A. Generating & Viewing the Reports

- In the task interface, under "**Preview responses**", you will see a link for each model's report. Click each link to open and view the pre-rendered HTML report.




- Open both reports:** Click the link for Report A and view it fully, then do the same for Report B.
- Take a first-pass look at both:** Before scoring anything, go through each report end-to-end. Get a general feel for the layout, visual quality, and how information is structured.
- Look for defects:** Note any broken elements, misaligned content, missing charts or tables, placeholder text, overlapping elements, or anything that appears incomplete or incorrectly rendered.

B. Cross-Checking with an External Browser

The built-in viewer **may not render every report perfectly**. If you want to verify how a report truly looks, you can download the HTML file and open it directly in a browser such as Chrome or Firefox.

To download the HTML file: **right-click the report link** → select **"Save Link As..."** → save it with a **".html"** extension → open the saved file in your browser.



- **Use this when something looks off:** If a chart, table, or layout element looks broken or unusual in the built-in view, download and open the file in a browser to confirm whether it is a viewer bug or a real issue in the report.
- **Resolve discrepancies through the code:** If the browser and the built-in viewer still show different results, inspect the HTML source code to determine whether it is a rendering compatibility issue or an actual defect in the report.



Important: If the built-in UI and the external browser show different results, always treat the HTML code as the ground truth.



NOTE: Do not assess factual accuracy in this step or anywhere in the evaluation. Assume all data points, facts, and metrics are correct. Your focus is entirely on design, layout, presentation, structure, and communication quality.



Step II: Understanding the Scoring Scale

Before assigning any scores, you must internalize the scoring rubric. Miscalibrated scores (particularly those that are lenient on critical issues) are one of the most common sources of misunderstandings in this project.

A. The Scoring Scale

All reports are scored on a **0–10 scale** across two pillars (Visual and Narrative). **Neither of the pillars has more weight; treat them both as equally important.** Use the table below as your anchor:

Score	Label	Definition
9–10	☀️ Exceptional	You would hand this to a CEO or print it in a consulting deck without changes.
7–8	💼 Professional	You would send this to a client or manager with minor edits, nothing structural.
5–6	📄 Average / Draft	You would circulate this internally for review before it goes anywhere near a client or leadership audience.
3–4	⚠️ Poor	You would send this back to the author for a significant rework before it leaves the team.
0–2	❌ Critical Failure	You would not circulate this at all: it is broken, incomplete, or fails to address the task.

B. The Standard to Keep in Mind

These reports are evaluated against a professional, publication-ready benchmark. Think: a management consulting deliverable, a data science briefing, or a slide deck prepared for

executive leadership. A report that *functions* but looks like an internal rough draft is **not** a 7 or 8: it is a 5 or 6.

Ask yourself: *Would I be comfortable presenting this report at a leadership meeting, unchanged?*

- If yes → the report is in the **7–10** range.
- If it needs polish → it belongs in the **5–6** range.
- If it has real structural or visual problems → **3–4**.
- If it is broken or near-unreadable → **0–2**.

C. Calibration Rules

Applying the scale correctly requires discipline. Follow these rules:

1. **Do not cluster scores in the 6–9 range by default.** If every report you score lands between 6 and 9, your calibration is too lenient. Reserve 9–10 for genuinely exceptional work. A 9 should be rare since it’s a report that is ready to be sent, no changes needed.
2. **A missing or broken element is never a 7+.** If a chart fails to render, a table is missing, or a section is incomplete, the report cannot score in the Professional range for that pillar; regardless of how good the rest of it looks. *Would you say it looks professional when a chart has some overlapping?*
3. **Weight issues by their impact on the reader.** Not all flaws are equal. A slightly off-color palette is a *minor* deduction. A misschoice of colors on a chart that prevents an effective, quick scan is a *major* issue. A wall of unformatted text, a broken layout, or missing data visualizations are serious failures. Let severity guide your score, not just the count of issues.

D. Issue Severity Reference

Use this as a quick guide when calibrating how much to penalize a given flaw:

Severity	Examples	Score Impact
 Critical	Broken layout, unrendered sections, missing tables/charts, illegible text	Can drop score to 0–4 depending on the extent
 Major	Text walls with no structure, poor data labeling, inconsistent formatting throughout	Typically keeps score at 5–6

 Moderate	Weak color palette, minor alignment issues, verbose prose	May reduce from 8 to 7, or 7 to 6
 Minor	Suboptimal hyperlink format, slightly uneven spacing, an issue in a chart that does not affect readability.	At most 1-point deduction; score floor of 8

Issues **do pile up**. A single minor flaw (a slightly off-color) may barely move the needle, leaving a report comfortably at an 8 or even a 9. But when **minor issues repeat across the report**, they stop being isolated mistakes and start revealing a pattern of low attention to detail. At that point, the score should drop accordingly, even if no single issue would have warranted it alone

E. A Note on Leniency

Leniency is the single most common calibration error in this project. It typically appears as:

- Giving a 7 to a report with broken elements "because the rest was good."
- Giving an 8 to a report "because it covered everything," ignoring that coverage is a Narrative criterion, not a Visual one.
- Avoiding low scores because the report "tried hard" or "was mostly there."

Your scores must reflect the output as a finished deliverable, **not the effort behind it or its potential**. Score what is on the page.



Step III: Visual Scoring

Score each report independently on **Pillar 1: Visual & Pixel-Level Aesthetics**. Provide a score and written rationale for each report, then write a side-by-side visual comparison.

A. Score Each Report Visually (Repeat for Report A and Report B)

- **Visual Score (0–10):** Rate the report's overall visual quality on a 0–10 scale.
- **Visual Score Rationale:** Briefly explain your visual score — what made it strong or weak?

Visual Score *

Visual Score (0–10).

Rate the report's overall visual quality. Consider:

- **Layout & alignment.** Are elements cleanly aligned and consistently spaced?
- **Color & typography.** Is the palette coherent and professional? Are fonts modern and readable?
- **Charts & tables.** Are they well-labeled, appropriately chosen, and free of clutter?
- **White space.** Does the layout breathe, or does it feel cramped?

⚠️ Do not consider content accuracy. Focus purely on how the report looks and whether it communicates information clearly through its design.

0

1

2

3

4

5

6

7

8

9

10

Updated ⓘ

Visual Score Justification *

 **Visual Rationale.** Briefly explain your visual score for this report.


Complete this for **Report A first**, then repeat the same process for **Report B**. Be rigorous, avoid clustering scores in the 6–9 range without clear justification.



Important: The individual justification for a particular report should avoid comparing features with those of the other report (e.g., if justifying Report A's visual score, do not add any comparison with Report B).

B. Write the Visual Side-by-Side Comparison

After scoring both reports individually, write a comparative analysis in the **"Visual side by side"** field. This should explain which report looks better overall and why, referencing specific observations from both.



C. Pillar 1 Scoring Criteria

Evaluate visual quality based on the four criteria below:

1. Professional Production Quality

- **Alignment & Grid:** Are elements (text, charts, logos) perfectly aligned? Does the image look "messy" or loosely structured?
- **Consistency:** Are font sizes, styles, colors, and bolding consistent throughout? Is spacing between elements consistent?
- **Comprehensiveness and Detail:** Does the artifact have the level of detail expected from an expert professional?
- **Completeness:** Are there missing elements (images, values, graphs) or placeholders suggesting a draft?

2. Aesthetics & Visual Hierarchy

- **Style:** Is the visual style adapted to the content domain? Are photos/logos appropriate for tone?
- **Scannable:** Is it easy to identify key information? Does the organization guide the user through the content?
- **Color Palette:** Does the artifact use a coherent, professional color scheme or default/clashing colors?
- **Typography:** Is the font modern and readable (e.g., clean Sans-Serif)? Is there a clear distinction between headings and body text?
- **White Space:** Does the layout breathe, or does it feel cramped and cluttered?

3. Data Presentation (Tables & Text)

- **Readability:** Can data be read instantly? Is there adequate row padding and contrast?
- **Titles:** Do tables and diagrams have clear titles and formatted header rows?

- **Formatting:** Are numbers formatted correctly? Are columns right-aligned for numbers and left-aligned for text?
- **Noise:** Does the document use unnecessary heavy borders or gridlines, or cleaner dividers?
- **Truncation:** Are tables split across pages? Are column widths appropriate for the content?

4. Visualization (Charts & Graphics)

- **Chart Selection:** Are charts appropriate for the data? (No 3D effects, no exploding pie charts.) Are they complete?
- **Proper Labeling:** Are charts properly labeled (axes, legends, titles)?
- **Data-Ink Ratio:** Are non-essential elements removed (redundant legends, excessive axis lines)?
- **Labelling:** Are data points directly labeled, or must the user reference a confusing legend?
- **Utility:** Are visual elements well-justified — adding value, not just decoration?

D. Visual Scoring Guidelines

Use the scale below to calibrate your visual scores. **Avoid clustering scores between 6–9** — examine each report closely and score rigorously.

Score	Label	Definition
9–10	Exceptional	Flawless execution. Visuals are publication-ready and perfectly aligned.
7–8	Professional	Strong execution. Minor aesthetic imperfections but a highly competent business document.
5–6	Average / Draft	Competent but flawed. Unpolished internal draft.
3–4	Poor	Significant issues: cluttered layout, confusing labels, weak visual flow.

0–2	Critical Failure	Unreadable formatting or broken layout elements.
-----	------------------	--



Step IV: Narrative Scoring

Score each report independently on **Pillar 2: Report Design, Structure & Storytelling**. The process mirrors Step II — score each report, write a rationale, then write a side-by-side comparison. The only difference is that you are now evaluating the narrative quality rather than the visual quality.

A. Score Each Report on Narrative (Repeat for Report A and Report B)

- **Narrative Score (0–10):** Rate the report's narrative and structural quality on a 0–10 scale.
- **Narrative Score Rationale:** Briefly explain your narrative score — what made it strong or weak?

Complete this for **Report A first**, then repeat for **Report B**. Apply the same rigor as in Step II — be specific and reference concrete evidence.



Important: The individual justification for a particular report should avoid comparing features with those of the other report (e.g., if justifying Report A's narrative score, do not add any comparison with Report B).

B. Write the Narrative Side-by-Side Comparison

After scoring both reports on narrative, write a comparative analysis in the **"Narrative side by side"** field. This should explain which report communicates its findings more effectively and why.

C. Pillar 2 Scoring Criteria

Evaluate narrative quality based on the six criteria below:

- 1. Structure:** Prefer reports that use structured Markdown (bullet points, callout blocks, tables, effective bold/italics) to convey insights quickly. High-quality reports prefer images and effective structures over dense prose.
- 2. Conciseness:** If two responses are equally comprehensive, the one that uses fewer words to reach the same level of insight is superior.
- 3. Comprehensiveness / Narrative Quality:** A comprehensive report fully addresses all aspects of the prompt, answers all explicit and implicit sub-questions, includes all crucial facts and caveats, and does not omit important context from the sources.
- 4. Coherence:** Visuals should blend coherently into the report and be appropriately titled or referenced in the text.
- 5. Precision:** Are the findings derived accurately from the data? Are units and metrics handled correctly?
- 6. The "Executive Test":** If a CEO had 60 seconds to scan this, which report looks most professional and yields the most insight even without a deep read?

D. Narrative Scoring Guidelines

Use the same 0–10 scale as visual scoring:

Score	Label	Definition
9–10	Exceptional	Crisp, ideally structured, and compelling. Fully addresses the prompt with precision.
7–8	Professional	Strong narrative. Minor verbosity, but highly competent business communication.
5–6	Average / Draft	Competent but flawed. Occasional text walls, some sub-questions unaddressed.
3–4	Poor	Weak narrative flow or significant gaps in addressing the prompt.

0–2	Critical Failure	Total failure to address the core prompt or incomprehensible structure.
-----	------------------	---



★ Step V: Overall Preference & Final Submission

After completing the visual and narrative scoring, provide an overall preference rating and complete the three final submission fields.

A. Overall Preference Rating

Select your overall preference on a 7-point scale ranging from "A is much better" to "B is much better."

The preference rank should be determined by taking into account these response characteristics: accuracy, instructions following, safety, helpfulness, and formatting/writing style.

A is much better No preference B is much better

☐ A
☐ A
☐ A
☐ -
☐ B
☐ B

Preference Justification *

Score	Meaning
1	Report A is much better than Report B
2	Report A is better than Report B
3	Report A is slightly better than Report B
4	Report A and Report B are about the same
5	Report B is slightly better than Report A
6	Report B is better than Report A

7

Report B is much better than Report A

B. Three Submission Fields



Overall Preference Rationale

Explain why you favored one report over the other overall. Reference both visual and narrative observations. This should be a detailed, evidence-based justification.



Summary

Write **one or two sentences** stating which report wins and the main reason why. This is a short, direct conclusion.



Feedback

Record any additional observations, edge cases, or issues worth flagging; for example, unusual rendering behavior, ambiguous prompts, or anything that made scoring difficult.

If there is no feedback, you can leave this section blank.

CRITICAL: All three fields in this step are required. Do not submit with empty fields. Provide substantive rationales — vague justifications like "A is better" are not sufficient.



Important: There are some tasks that the prompt is broken and shows error, if that happens you should continue evaluating the report by using other model's report and mention this in the Feedback section.



Important: There are some tasks that last paragraph of the model response is "Do not add image contents directly as base64 strings in [report.md](#)."



Important: If you see unrendered LaTeX equations or mermaid charts, please do not penalize the outputs for that.

Step VI: Common Errors

This section documents the most frequent mistakes observed during rating. Review these carefully before submitting; they represent patterns that consistently lead to inaccurate or poorly calibrated scores.

Error 1: Leniency in Model Scoring

Adherence to scoring guidelines must be absolute. We have observed numerous instances where contributors assign scores of 5 or higher despite critical visual failures, such as a missing table or an entire section failing to render.

Example of a critical visual failure:

2. Deep Dive: Current Rental Market Trends

*** BEIJING RENTAL TRAJECTORY SINGAPORE RENTAL TRAJECTORY (The "Summer Squeeze" Glut) (Post-Boom Stabilization & Divergence)

High Inventory Listings (+12.2%) Public HDB Demand Private Condo Supply | | ▼▼▼ Falling Effective Rents Modest Rent Growth Flat / Softening Rents (Landlord Concessions up to 1k RMB) (2% - 5% in 2025/2026) (0% - 3% Adjustment) ***

Error 2: Misclassifying Text Walls as a Visual Issue

A dense, unbroken block of text with no headers, bullet points, or whitespace is a **narrative flaw**; not a visual one. It reflects poor structure and communication design (Pillar 2), not a failure of layout or aesthetics (Pillar 1).

- **Do NOT penalize the visual score** for text walls unless there are accompanying visual problems; e.g., poor typography, zero line-spacing, or content running off the page.
- **DO penalize the narrative score** under Structure and Conciseness when the report reads as a wall of prose that buries key insights.

Error 3: Writing Generic Rationales

A rationale that says "**the report was well-structured**" or "**visuals were clean**" is not acceptable. Every rationale must be **specific, evidence-based, and criterion-by-criterion**.

For the **Narrative Rationale**, you are expected to address all **six Pillar 2 criteria individually**; Structure, Conciseness, Comprehensiveness, Grounding, Precision, and the Executive Test. For each one, state whether the model performed well or poorly and briefly explain why.

- **Correct:** "The report is well-organized with clearly labeled sections and consistent use of bullet points, making it easy to scan. However, several sections repeat the same finding in slightly different words, adding length without adding insight (GIVE DETAILED EXAMPLES HERE). Coverage is strong; all major sub-topics are addressed with

supporting data, and visuals are referenced in context rather than dropped in without explanation. The writing is precise with no obvious metric errors, and a busy reader could extract the key takeaways within a minute."

- **Incorrect:** "The narrative was good overall, with some room for improvement in clarity."

Error 4: Scoring ASCII Tables Incorrectly

Some models render data tables using plain ASCII characters (dashes, pipes) instead of proper HTML tables. This is a suboptimal design choice, but it is not an automatic failure. Use the following calibration:

Condition	Score Range	Rationale
ASCII table is clean, columns are aligned, and data is readable	7–8 (Professional)	Functional and legible, but not a polished deliverable. Minor visual penalty for using a non-native table format.
ASCII table has formatting issues — misaligned columns, broken characters, uneven spacing	5–6 (Average / Draft) or lower	The formatting failure compounds the choice of a weak table style. Score lower if other visual issues are also present.

Good ASCII example:

LOUD FATTY EXPANSION MATRIX				
Market	Launch Decision	SAM Size (\$M)		Strategic Focus
NEW YORK	STRONG LAUNCH	\$133.3M (39.7%)		Premium Pricing Power
MISSOURI	STRONG LAUNCH	\$79.3M (39.9%)		Capture Rapid Expansion
MARYLAND	STRONG LAUNCH	\$32.2M (23.1%)		Capitalize on Multi-Packs
ARIZONA	SELECTIVE LAUNCH	\$139.7M (69.1%)		Defend Margins / Velocity
WASHINGTON	HOLD / RE-EVALUATE	\$94.2M (44.5%)		Severe Price Compression

Bad ASCII example:

INFUSED PRE-ROLL UNIT VOLUME SHARE BY PRICE TIER				
State	Value (<\$10)	Core (\$10-\$20)	Premium (\$20-\$30)	Ultra-Premium (\$30+)
MD	86.0%	11.6%	2.2%	0.2%
AZ	43.9%	25.6%	24.4%	6.1%
NY	4.6%	33.4%	15.1%	46.9%

Rule: Judge ASCII tables on execution quality, not just format choice. A clean ASCII table is a 7–8; a broken one is a 5–6 or below.

Error 5: Uncentered Tables with a Loose Left Column should not be penalized

Previously we instructed to fail these issues but based on the new guidelines we **MUST NOT** penalize this issue.

Group	Headquarters	Key Brands	FY2024 Revenue	Revenue Change vs. 2023
LVMH	Paris, France	Louis Vuitton, Dior, Givenchy, Fendi, Celine, Tiffany & Co., Moët & Chandon	€79.2 bn	–8.2%
Hermès	Paris, France	Hermès	€15.2 bn	+15.0% (CER)
Kering	Paris, France	Gucci, Saint Laurent, Bottega Veneta, Balenciaga, Valentino	€17.2 bn	–12.2%
Richemont	Geneva, Switzerland	Cartier, Van Cleef & Arpels, IWC, Jaeger-LeCoultre	~€20.6 bn	~+4.0%
Burberry	London, United Kingdom	Burberry	£2.97 bn	–3.9%
Prada Group	Milan, Italy	Prada, Miu Miu, Car Shoe	€5.17 bn	+9.5%

Error 6: Cramped Data in Tables and Charts


Models sometimes pack too much data into tables or charts without adequate padding, row height, font sizing, or whitespace. The result is a report that is technically complete but visually overwhelming and hard to read at a glance.


- **In tables:** Look for cells with no internal padding, rows that are too tight, text that runs to the edge of the cell, or columns that are too narrow for their content.
- **In charts:** Look for axis labels that overlap, legends that crowd the chart area, data points that are indistinguishable, or bars/lines that are too compressed to read.
- **Where to score it:** Penalize under Readability and Truncation (Data Presentation, Pillar 1) and White Space (Aesthetics, Pillar 1).

Example 1: overlapping text with shape



Example of cramped data in chart:





Example: The note is not centered below the table and break the formatting

BOM TARGETS & PRICING ROADMAP			
SKU Format	Target BOM	Wholesale Price	Proposed Retail MSRP
1.0g Single-Pack	\$1.47 + WIP	\$4.00 - \$10.00	\$12.00 - \$20.00
0.5g 5-Pack (2.5g)	\$3.07 + WIP	\$9.00 - \$22.50	\$30.00 - \$45.00

Note: Bill of Materials (BOM) targets include unbleached cones (\$0.40), specialized IML packaging (\$1.00), and master merch packs (\$0.07). Work in Progress (WIP) accounts for active terpene, rosin, THCa inputs, and CMO production labor.

Rule: A report can be data-rich and still breathe. Cramped presentation is a deliberate design failure, not just a side effect of having a lot of content.

Error 7: Markdown Rendering Failures

Markdown rendering failures are **visual defects** that should be reflected in the Pillar 1 score:

Issue	What It Looks Like	Scoring Impact
Markdown not rendered in bullet points	Bullet items display raw Markdown syntax, e.g., **bold** or - item instead of formatted text.	Penalize under Consistency and Completeness. If widespread, this signals a fundamental rendering failure and should significantly lower the visual score.

Example: table has incorrect reference numbers

Table of Contents

1. [Executive Summary](#)
2. [Introduction: The Digital Inflection Point](#)
3. [The European Luxury Landscape: An Overview](#)
4. [E-Commerce Growth and Market Dynamics](#)
5. [Brand-by-Brand Analysis](#)
6. 5.1 [LVMH — Controlled Scale and Omnichannel Mastery](#)
7. 5.2 [Hermès — The Art of Digital Resistance](#)
8. 5.3 [Kering / Gucci — A Cautionary Tale of Overexposure](#)
9. 5.4 [Burberry — Heritage Reloaded](#)
10. 5.5 [Prada Group — Minimalism Meets Metaverse](#)
11. 5.6 [Richemont — Maisons in the Digital Age](#)
12. [Digital Marketing Campaigns: Case Studies](#)
13. [The Exclusivity Paradox: Does Digital Dilute Luxury?](#)
14. [Consumer Behaviour and Generational Shifts](#)
15. [Financial Performance and Digital Investment Correlation](#)
16. [Emerging Technologies: AI, AR, Blockchain, and the Metaverse](#)
17. [The Collapse of Multi-Brand Luxury Platforms](#)
18. [Strategic Frameworks and Recommendations](#)
19. [Conclusion](#)
20. [References](#)

Example: Incorrect bullet point formatting

Your product architecture aligns perfectly with three high-value consumer personas: * **The Cannasseur (Relaxation & Recovery):** Experienced craft users who value evening rituals. They seek high THC potency paired with a clean, unblemished burn. Loud Fatty allows them to achieve their desired state efficiently with a single 1g joint rather than consuming multiple standard pre-rolls. * **The Trendsetter (Social Group Sessions):** Modern consumers who want sleek, elegant products for concerts, festivals, and parties. Loud Fatty's kief-free exterior allows seamless passing among friends without leaving residue on fingers, ensuring the host supplies the "loudest" and most stylish joint in the room. * **The Premium Buyer (Special Occasions & Gifting):** Consumers who purchase cannabis as a luxury indulgence or celebratory gift. Your translucent flip-top packaging and elevated extract formulation provide a top-shelf unboxing experience analogous to popping a fine bottle of champagne.

Error 8: Suboptimal Hyperlink Integration

Hyperlinks should be embedded directly into the relevant reference text rather than appended as standalone URLs without context.

Good Example:

- Business Tats. (2025). [Nike's Sales Worldwide Since 2005](#).

Bad Example:

- Business Tats. (2025). *Nike's Sales Worldwide Since 2005*.

<https://businesstats.com/nikes-sales-worldwide-since-2005/>

This should not be penalized harshly. This issue alone could only reduce the model's score to 8 at most.

Error 9: Comparison with other reports in the per-response justification

Any individual justification (narrative or visual) should only flag issues for that report without comparing to the other report.

Example (comparison in individual report justification):

The screenshot displays a user interface for providing a narrative score justification. On the left, there is a table with two rows: 'Response B' and 'Report B Here'. The 'Response B' row is highlighted in light blue. To the right of the table is a 'Narrative Score Justification' section. At the top of this section is a horizontal row of buttons numbered 0 through 10, with the number 7 selected and highlighted in blue. Below the buttons, the text reads: 'Narrative Rationale. Briefly explain your narrative score for this report.' A text box contains the following justification: 'Report B is a strong and highly comprehensive narrative in the professional band, a notch below Report A on writing and structure.'