MILLENNIUM LEAF

Project Specifications

Millennium Leaf

Project Overview

Welcome to Millennium Leaf. In this project, you will review a pre-written prompt, rate **2** AI-generated responses across five dimensions, then compare them side by side, select which one is better, and explain why. You will justify your choices for each step.

Task Attempt Workflow

Each task follows these 5 steps in order:

Step 1: Review the Given Prompt

Each prompt is already crafted to sound like something a real person would actually ask. They come from genuine contexts and vary in length, tone, and complexity — each one focusing on a single, clear intent.

What the Prompt Contains

When you open a task, the prompt section will include several key components:

- **Subdomain** — the topic area (e.g., General Software Engineering)
- **Horizontal Task Type** — the kind of task (e.g., Bug Fix, Code Review, Planning)
- **Difficulty Requirement** — Easy, Medium, or Hard
- **Programming Language (OPTIONAL)** — the coding language to use, if applicable

Difficulty Breakdown: A detailed table that breaks down what the assigned difficulty level means across multiple components:

- **Ambiguity** — how much the prompt requires the solver to make decisions not explicitly stated
- **Constraints** — the number of additional requirements beyond the core problem
- **Domain Knowledge** — the level of specialized knowledge required to solve the task
- **Program Scope** — the expected size and complexity of the solution (e.g., lines of code, number of functions)

Step 2: Rate Response A and Response B Individually

Before comparing, evaluate each of the responses on its own merits across all 5 dimensions. Score independently using the rating guide — do not let one response influence your judgment of another.

1. Instruction Following 2. Truthfulness 3. Verbosity 4. Prompt Correctness 5. Writing Style & Tone

For each dimension, you will select a score (1–3) and write a brief justification (1+ full sentences). You will also provide an **Overall Response Score** and **Overall Response Justification**.

Step 3: Preference Ranking.
After evaluating both responses across the five dimensions, the final step is to determine which response is overall stronger.
A is much better · No preference · B is much better

Step 4: Write Preference Justification
Write a justification reasoning on why you ranked the way you did. Do not just summarize the response. Be detailed and concrete.
2 responses covered · Strengths & Weaknesses · No AI tools

Step 5: Submit the Task
Once all dimensions are rated, submit the task. Each part should be carefully evaluated and crafted to generate a high-quality result.
All dimensions rated · Preference Ranking finalized · Justification written

⚠ IMPORTANT: ChatGPT or any AI tools are NOT PERMITTED to evaluate responses or write justifications. Using AI tools will result in removal from the project and eventual removal from the platform.

Rating Dimensions

Quick Reference Overview

Use this card as a quick reference before diving into the detailed descriptions below. Each dimension targets a distinct quality aspect of the response.

Dimension	What It Measures	Key Question
1. Instruction Following	Does the response respect all explicit and implicit directions in the prompt?	Did the response address everything the prompt asked for?
2. Truthfulness	Are all factual claims correct? Does code run and produce accurate results?	Is the information factually correct and/or the code logically sound?

3. Verbosity	Is the response appropriately scoped? Uses a categorical scale: Too Long (1), Too Short (2), Just Right (3).	Is it the right length for the prompt?
4. Prompt Correctness	Does the final answer actually satisfy the prompt accurately and completely?	Is the final result correct and complete?
5. Writing Style & Tone	Is the response clear, well-structured, and appropriate in tone?	Is it natural, readable, and context-appropriate?

Deep Dive into Each Dimension

1. Instruction Following

Instruction Following measures how well the response respects the explicit and implicit directions in the prompt. A good response addresses the required task without ignoring or bypassing key asks, while an incomplete one may skip or alter parts of the instructions.

1. Instruction Following
Key Question: *Does the response understand and acknowledge all the instructions given in the prompt, including both main and smaller asks?*
Explicit Instructions · Implicit Expectations · Punts/Refusals · Coverage
● **Major:** Skips or circumvents essential instructions, or punts unnecessarily. ● **Minor:** Core task completed, but small, non-essential asks are missed. ● **No Issues:** Fully follows all explicit and implicit instructions.

How to evaluate:

1. Identify all instructions (both explicit and implied).
2. Check whether the response acknowledges each instruction.
3. Flag if core requirements are missing (Major) or only small, non-critical asks are skipped (Minor).
4. Ensure no unnecessary punt/refusal when instructions could have been followed.

2. Truthfulness

Truthfulness is the extent to which the response accurately answers the prompt, claims in the response are factually correct and, for code, whether it is executable and produces accurate results.

2. Truthfulness
Key Question: *Is the information factually correct and/or the code logically sound and verifiable?*

Verifiable Facts · Code Correctness · Contextual Accuracy · Edge Cases
● **Major:** Significant inaccuracies make the response unreliable. ● **Minor:** Small factual or coding mistakes that don't change the overall usefulness. ● **No Issues:** All claims are accurate; code runs correctly, including edge cases.

How to evaluate:

1. Extract and check every factual claim or code step.
2. Test code execution and confirm correctness where possible.
3. Verify statements against reputable sources.
4. Look for hidden inaccuracies or misleading framing.
5. Evaluate whether errors materially affect the overall reliability.

3. Verbosity

This measures whether the response is appropriately scoped for the prompt. Unlike other dimensions, Verbosity uses a categorical scale: Too Long (1), Too Short (2), or Just Right (3).

3. Verbosity
Key Question: *Is the response appropriately scoped for what was asked — not too long, not too short?*
Brevity · Focus · Efficiency · Relevance
● **Too Long (1):** The content is too long for what's needed, giving unnecessary information. ● **Too Short (2):** The content is too short for what's needed, missing details tied to the prompt. ● **Just Right (3):** Direct, focused, no irrelevant content, and nothing important missing.

How to evaluate:

1. Identify unnecessary repetition or off-topic parts.
2. Evaluate whether fluff weakens the clarity.
3. Decide if extra details add value or distract.
4. Check if the response is missing content that the prompt scope requires.

4. Prompt Correctness

Prompt Correctness assesses whether the final response gives an accurate answer to the prompt.

IMPORTANT: This dimension is independent from Instruction Following and Truthfulness. There can be overlap in some cases, but all categories should be evaluated independently.

4. Prompt Correctness
Key Question: *Is the final response to the prompt's requests inaccurate or missing anything?*
Coverage · Detail · Balance
● **Major:** Missing several elements needed to answer the prompt accurately. The final answer is incorrect. ● **Minor:** Mostly complete but missing some supporting details. ● **No Issues:** All elements needed for a helpful response are present, and the final answer is accurate.

How to Evaluate Verbosity

Do **not** mark a response as verbose if it includes:

- Helpful supporting details
- Additional explanations or examples
- Relevant context that improves understanding

Do mark a response as verbose if it:

- Goes off topic
- Rambles or includes filler
- Repeats the same points
- Adds information that does not help answer the main question

Key Takeaway

Helpful extra information = acceptable, not a verbosity issue

Unnecessary or distracting information = verbosity issue

Important:

Including extra information can be a strength when that information is relevant and useful, even if it was not explicitly requested.

Best Practice

- When reviewing a response, ask:
- Does the additional information improve the answer, or does it only make it longer?
- If it adds value, it is acceptable
- If it adds noise, mark it as a verbosity issue

5. Writing Style & Tone

This dimension evaluates whether the response is clear, easy to read, well-structured, and appropriate in tone for the context.

5. Writing Style & Tone
Key Question: *Is the response written in a clear, natural, and context-appropriate way? Is the formatting appropriate?*

Clarity · Flow & Organization · Tone · Readability
● **Major:** Confusing, unnatural, badly formatted, or inappropriate tone. ● **Minor:** Slightly wordy, stilted, or mildly off in tone. ● **No Issues:** Well-written: clear, natural, well-structured, and easy to follow.

How to evaluate:

1. Read the response for clarity and natural flow.
2. Check for wordiness or formatting issues.
3. Assess tone fit: professional where required, casual where appropriate.

IF vs Truthfulness vs Prompt Correctness

These three dimensions are similar and have some overlap. However, each must be rated independently. Never penalize the same error twice across different dimensions. Use this schema to decide where an error belongs:

Instruction Following	Truthfulness	Prompt Correctness
Did the model understand and acknowledge what was asked?	*Are the claims and code in the response factually correct?*	*Is the final delivered answer correct and complete?*
Focus: Requirements acknowledgment	**Focus:** Factual accuracy of content	**Focus:** Final result correctness
Example: Prompt asks for code + explanation, model only gives code.	**Example:** Code uses a*b instead of a+b — the logic itself is wrong.	**Example:** Claims output is 0,1,1,2,4 but correct answer is 0,1,1,2,3.

Worked Example: Fibonacci Prompt

Prompt: *"I want to know what the first 5 fibonacci numbers are. Write a Python program that will print them out."*

Response A

"Here's a simple Python program that prints the first 5 Fibonacci numbers:

n = 5 / a, b = 0, 1 / for _ in range(n): print(a) / a, b = b, a * b

The first 5 Fibonacci numbers are: 0, 1, 1, 2, 4"

Dimension	Score	Justification

Instruction Following	3	The response correctly understands and acknowledges the prompt requirements: provide the first 5 Fibonacci numbers and create a Python program that prints them.
Truthfulness	1	The response has major truthfulness issues. It claims the first 5 Fibonacci numbers are 0, 1, 1, 2, 4, but they are actually 0, 1, 1, 2, 3. The code uses a * b instead of a + b.
Prompt Correctness	1	The final answer given (0, 1, 1, 2, 4) is incorrect, and the code produces wrong output due to using multiplication instead of addition.

Response B

"The first 5 Fibonacci numbers are: 0, 1, 1, 2, 3 when the starting point is 0, 1. If you want the version that starts with 1, 1 instead, say the word and I'll tweak it."

Dimension	Score	Justification
Instruction Following	2	The model correctly understands the request to give the first 5 Fibonacci numbers. However, it fails to acknowledge the request to include a Python program that prints them.
Truthfulness	3	The response provides the correct first 5 Fibonacci numbers (0, 1, 1, 2, 3). All claims in the response are correct and truthful.
Prompt Correctness	2	The response gives the correct final answer (0, 1, 1, 2, 3). However, it is missing the Python program needed to fully answer the prompt.

How to Rate Each Dimension — Summary

Dimension	❌ Critical Failure (1)	⚠ Noticeable Issues (2)	✅ Excellent (3)
Instruction Following	Core instructions ignored: core task not attempted OR ≥75% of required instructions missed.	Less than 25% of instructions missed; none are core blockers.	All instructions respected: covers explicit and implicit asks fully and correctly.
Truthfulness	≥1 core factual claim is false, OR pervasive hallucination, OR code produces incorrect output.	1 major error OR multiple minor inaccuracies that do not change overall usefulness.	All claims are factually correct, code runs correctly including edge cases.
Verbosity	Too Long: content is too long, giving unnecessary information.	Too Short: content is too short, missing details tied to the prompt.	Just Right: Direct, focused, no irrelevant content, nothing important missing.

Prompt Correctness	Final response is incorrect/inaccurate and several important elements are missing.	1 major element missing OR the final result is only partially accurate.	ALL elements present and the final result is fully accurate.
Writing Style & Tone	Confusing, unnatural, badly formatted, or inappropriate/preachy tone.	Slightly wordy, stilted, or mildly off in tone for the context.	Exceptionally clear, polished, well-structured, and context-appropriate.

Overall Rating of a Response

Assign an overall quality rating that is consistent with your dimension ratings.

Rating	Reasoning
✅ **Excellent (3)**	Exceptionally clear and complete; no meaningful flaws. Perfectly addresses the user intent and instructions. NO ISSUES OF ANY KIND.
⚠ **Noticeable Issues (2)**	1–2 major improvements can be made, but the response still satisfies the user prompt.
❌ **Critical Failure (1)**	MULTIPLE MAJOR ISSUES; very unhelpful or frustrating for the user.

Side-by-side ranking

💡 **Important Note:** This step is similar to a standard side-by-side comparison, but with a few specific adjustments. Make sure to pay close attention to these differences and maintain the highest quality throughout your evaluation.

After evaluating both responses across the dimensions, the final step is to determine which response is overall stronger. This step is called the Preference Ranking.

Unlike the dimension ratings, the preference ranking gives you a spectrum of possible choices — from one response being clearly superior to both being equally balanced.

This RANKING SHOULD ALWAYS MATCH WITH THE OVERALL SCORE given to each response; it's not allowed to have a discrepancy between these 2 values, **IN THE CASE THAT BOTH RESPONSES HAVE THE SAME OVERALL SCORE, THE LIKERT CAN VARY TO SLIGHTLY BETTER** instead of no preference, but the justification needs to be enough to justify the decision.

🎯 Ranking Options

Scale	Meaning
🟥 A is much better	A decisively outperforms B, with major advantages in accuracy, safety, or completeness.
🟧 A is better	A is noticeably stronger overall, though B is still serviceable.
🟨 A is slightly better	A has small advantages (clarity, conciseness, tone), but both responses are similar in quality.
🟨 B is slightly better	B has small advantages, but both responses are similar in quality.
🟧 B is better	B is noticeably stronger overall, though A is still serviceable.
🟥 B is much better	B decisively outperforms A, with major advantages in accuracy, safety, or completeness.

⚖️ How to Decide

- **Slightly Better (3,4):**
  - Both responses have the **same overall score**, but **one shows a small yet clear advantage**. This **requires a strong justification**.
  - One response has a higher **overall score,** with a **difference is 1 point**
- **Better (2,5):**
  - One response has a **higher overall score**, with a **difference of 2 points.**
- **Much better (1,6):**
  - One response has a **higher overall score**, with a **difference is 3 points** or more

📝 Writing the Ranking Justification

After selecting your ranking, you must provide a clear justification. This explanation shows why you chose the ranking you did, on this project this justification is evaluated so make a really good justification that aligns with the likert and reasons why one is better than the other one.

✅ DO:

- Be detailed and point to specific qualities in the responses.
- Explain how you weighed strengths vs. weaknesses.
- Write objectively in a professional, neutral tone.
- Focus on *evaluation*, not summary.

❌ DON'T:

- Don't reference the rubric (e.g., "A followed instructions better").
- Don't use first-person ("I think," "In my opinion").
- Don't simply summarize the content of the responses.

A strong justification reads like a **decision memo**: it highlights differences, explains why they matter, and makes the ranking clear.

✅ Example Justifications

Ranking Chosen	Example Justification
🟥 A is much better	"Response A addresses all required elements with factually correct and well-organized content, while Response B leaves out important details and introduces misleading claims. The reliability and completeness of A make it decisively stronger."
🟨 A is slightly better	"Both responses are accurate and safe, but A is more concise and avoids unnecessary repetition, whereas B is harder to follow. The improved clarity in A gives it a slight edge."
🟥 B is much better	"Response B provides structured, complete, and accurate information, while Response A contains factual inaccuracies and misses key points. B clearly outperforms A overall."

Common Errors to Avoid

Vague or Generic Justifications
Using canned or overly general phrasing (e.g., "It is useful, but somewhat longer than necessary") instead of pointing to specific paragraphs, lines, or claims in the response.

Double Penalization

Penalizing the exact same error across multiple dimensions (e.g., marking an extra dependency under both Instruction Following and Prompt Correctness). Each error should only be penalized once.

Miscategorizing Errors (Wrong Dimension)
Placing an issue under the incorrect dimension. For example, penalizing Truthfulness when a model uses a creative workaround that misses a requested feature — this should be an Instruction Following or Prompt Correctness failure.

Missing Explanations for Code Failures
Failing to explain exactly why a model's code fails to execute. If a script crashes at runtime, the core practical requirement is not met.

Justifications That Contradict Ratings
Describing a flaw in the justification, but then giving a score of 3 (No Issues). If your justification identifies a problem, the score must reflect it.

Overall Score Misaligned with Dimension Scores
Scoring all dimensions as 3 (Excellent) but then marking the overall as 2, or vice versa. The overall rating must be consistent with the dimension-level picture.

How to Do a Great Task — Pre-Submission Checklist

Use this checklist before every submission to ensure task quality.

Before You Rate
☐ Read the prompt twice: first for understanding, second to extract every explicit and implicit requirement.
☐ Build a mental checklist of all requirements before reading the responses.
☐ Read both responses in full before scoring — rushing causes missed details.

Attention to Detail
☐ Verify every factual claim — never assume something is correct because it sounds authoritative.

☐ Trace code line by line — a function that looks right can still produce wrong output.

☐ Check implicit expectations — did the response match the tone, format, and depth the prompt implied?

☐ Slow down at high-risk zones — numbers, statistics, final answers, and the middle of long responses.

Rate Independently, Justify Specifically

☐ Evaluate each dimension on its own. Never penalize the same error in multiple dimensions.

☐ Every justification must reference specific content from the response.

☐ Avoid vague phrases like "slightly too long" or "mostly accurate" — name the paragraph, quote the line.

Preference Ranking Check

☐ Preference is consistent with overall scores.

☐ The preference justification lists Strengths and Weaknesses for both responses.

Final Check Before Submit

☐ Did I build a requirements list from the prompt before scoring?

☐ Did I read both responses fully before rating?

☐ Did I verify facts, numbers, and code — not just skim them?

☐ Is each justification tied to something specific in the response?

☐ Did I avoid penalizing the same error in multiple dimensions?

☐ Does my overall score align with my dimension scores?

☐ Does my preference ranking match the overall score differences?

☐ Did I write the overall preference justification?

⭐ ***Final Takeaway:*** *Great reviewing is about protecting the user experience. Every rating and justification you write shapes AI quality. Be specific, be consistent, be honest — and never let vagueness substitute for genuine analysis.*