

← Back to Outlier

Openclaw Atlas

Intro

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Claw & Order

How to Rubric



Outlier

MAIN CAPABILITIES - RUBRIC
SCREENING

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

What are rubrics?



Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

What are Rubrics in OpenClaw?

Definition

A rubric is a list of **Present / Not Present criteria** used to evaluate whether a model successfully completed a task.



What Rubrics Evaluate

- Agent behavior and decision-making
- Reasoning quality
- Tool usage
- Final output quality



Important Notes

- I. Rubrics must be based on the **specific task and desired outcome**.
- II. **Avoid generic criteria** that could apply to any task.

A good rubric measures observable results, not opinions.

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Outlier

Outcome-First
Philosophy

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min



Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min



Almost Right: The Science of
LLM Errors
5 sections • 62 min



OpenClaw RL Common Errors
Screening
2 sections • 31 min



 Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min



OpenClaw Agentic Task
Simulator



Outcome Criteria Only

Rubrics should evaluate whether the model successfully achieved the desired outcome, not simply whether it demonstrated good reasoning.

Before writing a rubric, ask yourself:

- ? Did the model complete the main objective?
- ? Did it produce the required output or artifact?
- ? Did it follow all required formats, rules, and constraints?
- ? Can the result be clearly verified as correct or incorrect?

Outcome-First Evaluation

The final outcome is the primary measure of success. While reasoning may support the result, rubrics should focus on what the model actually produced. If the desired outcome is unclear or difficult to verify, creating effective rubrics becomes much harder.

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Rules



Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Each Criterion Must Be Self-Contained

A criterion is not self-contained if you must look at the prompt, reference materials, another criterion, or outside information to evaluate it.

Ask yourself:

If I only had the model response, could I determine whether this criterion passed or failed?

→ If the answer is no, the criterion is not self-contained.

Example 1

✗ Not Self-Contained

"Response identifies the first president of the USA."

✓ Self-Contained

"Response identifies George Washington as the first president of the United States."

Example 2

✗ Not Self-Contained

"The response addresses the bug mentioned in the prompt."

✓ Self-Contained

"The response addresses the issue where the submit button does not function."

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Each Criterion Must Be Atomic

A criterion should evaluate one requirement only.

If a criterion checks multiple requirements at the same time, it becomes difficult to determine exactly what passed or failed.

Ask yourself:

Can this criterion be answered with a single clear Yes or No?

→ **If not, it likely needs to be split into separate criteria.**

Example

✗ Not Atomic

"The agent includes columns named 'party,' 'season,' and 'beverages.'"

If one column is missing, should the criterion pass or fail?

✓ Atomic

- The agent includes a column named party.
- The agent includes a column named season.
- The agent includes a column named beverages.

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Each Criterion Must Be Objective

A criterion should be based on clear, measurable facts, not personal judgment or interpretation.

If two reviewers could reasonably disagree on whether a criterion passes or fails, the criterion is probably not objective.

Ask yourself:

Can this criterion be evaluated with a clear Yes or No answer?

→ If not, make it more specific.

Example

✗ Not Objective

"The response should have good formatting."

What does "good" mean?

✓ Objective

- "The response includes a title."
- "The response uses bullet points."
- "The response contains a table with three columns."

● Avoid vague words such as:

- Good
- Appropriate
- Reasonable
- Helpful
- Clear
- Well-written

Unless they are supported by specific, measurable requirements.

Part 1: Onboarding ^

OpenClaw - Screening Course ✓
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes 🔒
3 sections • 1 min

Almost Right: The Science of LLM Errors 🔒
5 sections • 62 min

OpenClaw RL Common Errors Screening 🔒
2 sections • 31 min

🔍 Task Quality Module:
Rubrics, Checks & Trajectories 🔒
6 sections • 46 min

OpenClaw Agentic Task Simulator 🔒

Each Criterion Must Have a Category

Every rubric criterion should be assigned to the category that best matches what it evaluates.

Task Completion

Evaluates whether the model successfully completed the requested task and produced the required outcome.

Instruction Following

Evaluates whether the model followed specific requirements, constraints, formats, or rules.

Factuality & Hallucination

Evaluates whether the response is accurate and supported by available information, without adding unsupported claims.

Tool Use

Evaluates whether the model used the appropriate tools correctly and effectively to complete the task.

Agent Behavior

Evaluates the model's decision-making, planning, memory usage, and overall agent behavior throughout the task.

Important Note

If a criterion could fit multiple categories, select the category that best represents its primary purpose.

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Criteria



Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Ensuring Complete Rubric Coverage

Keep Rubrics Concise and Focused

- Include only the most important checks needed to evaluate success.
- Avoid duplicate, overlapping, or unnecessary criteria.

Cover Critical Events

- Every critical step required to achieve the Desired Outcome should have at least one rubric criterion.
- If a step can fail independently, it should be evaluated independently.

Examples of Critical Events

- ✦ Producing the required final artifact or file.
- ✦ Completing required integrations or tool actions.
- ✦ Following key decision rules or scoring requirements.
- ✦ Generating the required output format or structure.

Important Reminders

Do not combine separate failure points into a single criterion.

If two requirements can fail independently, create separate criteria for each one.

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Outlier

15

Handling
Repeated Actions

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Handling Repeated Actions with Spot Checks

When a task contains many repeated actions (emails, rows, records, entries, etc.):

- Do not create a separate rubric criterion for every item.
- Instead, use a combination of count verification and spot checks.

Count Verification

Confirm that the correct number of actions were completed.

Example

- The agent sends all 16 required emails.
- The agent creates all 25 required records.

Spot Checks

Review a small sample of items in detail to verify quality and accuracy.

Example

- The agent correctly sends email #2.
- The agent correctly sends email #4.
- The agent correctly sends email #7.

Why Use Spot Checks?

- ❖ Keeps rubrics concise and manageable.
- ❖ Prevents hundreds of nearly identical criteria.
- ❖ Still provides confidence that the task was completed correctly.

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min


Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Weights



Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Positive Wording & Negative Criteria



Positive Wording Requirement

All rubric criteria should be written as clear, direct statements. The point value determines whether the criterion rewards or penalizes the response.



Avoid Negative or Double-Negative Wording:

Do not write criteria like:

- ✗ "The response does not fail to..."
- ✗ "The response does not incorrectly..."
- ✗ "The response does not omit..."

These are difficult to evaluate and can create confusion.

Good Example of a Negative Criterion:

"The response includes French words." (-3 points)- (The task required English-only words.)

Notice that the criterion is written as a direct statement. The negative score is what penalizes the response.

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min



Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min



Almost Right: The Science of
LLM Errors
5 sections • 62 min



OpenClaw RL Common Errors
Screening
2 sections • 31 min



 Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min



OpenClaw Agentic Task
Simulator



Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Positive Weights

Every rubric criterion must use one of the approved positive weights:

+5 Critical

- Required for the task to be successful.
- Directly measures whether the core objective was achieved.
- If this criterion fails, the task should generally be considered unsuccessful.

Examples

- ❖ The required file is created.
- ❖ The final report contains all required sections.
- ❖ The agent completes the primary task requested.

+3 Important

- Measures an important requirement that strongly impacts quality.
- Failure does not necessarily make the entire task unsuccessful, but it significantly reduces quality.

Examples

- ❖ The output follows the required format.
- ❖ The agent uses the required data source.
- ❖ Key instructions are followed correctly.

+1 Nice to Have

- Improves quality, clarity, or usability.
- Failure has little impact on overall task success.

Examples

- ❖ The output is well organized.
- ❖ The response includes helpful formatting.
- ❖ The final artifact is easy to review.

Quick Rule:

+5 = Required for success

+3 = Important for quality

+1 = Helpful improvement

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Negative Weights

Every rubric set must include at least one negative criterion.

-1 Minor Issue

- Small mistake that has limited impact on the final result.
- The task is still largely correct and usable.

Examples

- ❖ Minor formatting issue.
- ❖ Small inconsistency in wording.
- ❖ Extra information that does not affect correctness.

-3 Major Issue

- Significant problem that reduces quality, accuracy, or reliability.
- The task may still be usable, but an important requirement was not met.

Examples

- ❖ Missing an important instruction.
- ❖ Incorrect formatting of a required output.
- ❖ Incomplete execution of a key step.

-5 Critical Failure

- Serious issue that makes the result incorrect, unusable, or invalid.
- Violates a required constraint or prevents successful task completion.

Examples

- ❖ Hallucinated information or tool results.
- ❖ Missing the required final artifact.
- ❖ Ignoring a critical instruction.
- ❖ Taking an action that should not have been taken.

Quick Rule:

-1 = Small issue

-3 = Major issue

-5 = Critical failure

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course ✓
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes 🔒
3 sections • 1 min

Almost Right: The Science of LLM Errors 🔒
5 sections • 62 min

OpenClaw RL Common Errors Screening 🔒
2 sections • 31 min

🔍 Task Quality Module: Rubrics, Checks & Trajectories 🔒
6 sections • 46 min

OpenClaw Agentic Task Simulator 🔒



Important Note

Only these weights are allowed:

-5, -3, -1, +1, +3, +5

Using any other value makes the rubric invalid.

Outlier

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Common Weighting Mistakes

Always make sure each criterion has the correct weight. Incorrect weights can unfairly increase or decrease a model's score.



Major Weighting Error

A criterion is assigned a weight that is two levels away from the correct value.

Examples:

- ★ Using **+1** when **+5** is appropriate
- ★ Using **-1** when **-5** is appropriate
- ★ Using **+5** when **+1** is appropriate

These mistakes significantly affect the final score and should be considered major issues.



Minor Weighting Error

A criterion is assigned a weight that is one level away from the correct value.

Examples:

- ★ Using **+1** instead of **+3**
- ★ Using **+3** instead of **+5**
- ★ Using **-1** instead of **-3**
- ★ Using **-3** instead of **-5**

These mistakes still affect scoring but have a smaller impact.



Quick Check

Before submitting your rubric, ask yourself:

- ? Is this criterion slightly important (**±1**)?
- ? Is it important (**±3**)?
- ? Is it critical to task success (**±5**)?

Choosing the correct weight helps ensure the final score accurately reflects the model's performance.

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min



Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min



Almost Right: The Science of
LLM Errors
5 sections • 62 min



OpenClaw RL Common Errors
Screening
2 sections • 31 min



 Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min



OpenClaw Agentic Task
Simulator



Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Scoring Rubrics

For each rubric criterion, select one of two options:

PRESENT (Meets the Criterion)

- The model successfully satisfies the rubric requirement.
- The expected behavior, action, or output is present.

NOT PRESENT (Does Not Meet the Criterion)

- One or more required parts of the criterion are missing.
- The requirement was not fully completed.

Common Reasons a Criterion is NOT PRESENT

- ✦ The required artifact was not created.
- ✦ A required instruction was not followed.
- ✦ A constraint was violated.
- ✦ The output format is incorrect or incomplete.
- ✦ A required action was skipped.

Important Reminders

Not every model is expected to satisfy every criterion.

A strong rubric should help distinguish performance differences between models. If every model passes nearly every criterion, the rubric may not be challenging enough.

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator



Scoring Note

How scoring works:

- If a positive criterion is **PRESENT**, its weight is added to the model's score.
- If a negative criterion is **PRESENT**, its weight is applied as a penalty.
- All criteria are evaluated using the same **PRESENT / NOT PRESENT** process.

Outlier

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Ratings

Each rubric criterion should be marked as either **PRESENT** or **NOT PRESENT**.

PRESENT

The criterion is fully satisfied.

- ✓ All required parts of the criterion are met.
- ✓ The model's output matches the rubric requirement.
- ✓ No important elements are missing.

What to write:

- Briefly describe what was observed and why the criterion was met.
- The source must be relevant to the task objective.

Not Allowed

The criterion is not fully satisfied.

- ✗ One or more required parts are missing.
- ✗ The output does not meet the rubric requirement.
- ✗ A required action, format, or constraint was not followed.

What to write:

- Briefly explain what is missing or why the criterion was not met.

Score carefully
Incorrectly marking a criterion as PRESENT or NOT PRESENT can significantly affect the final score and ranking of a model.

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min



Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min



Almost Right: The Science of
LLM Errors
5 sections • 62 min



OpenClaw RL Common Errors
Screening
2 sections • 31 min



 Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min



OpenClaw Agentic Task
Simulator



Task example



Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Example Desired Outcome



- An Anki import file is created in the workspace (e.g., spanish101_anki_import.tsv). The file is a TSV file containing 14 data rows.
- The first column contains the Spanish word, and the second column contains the English translation.
- Flashcards are created for exactly 14 words:
 - casa → house
 - mesa → table
 - libro → book
 - escuela → school
 - puerta → door
 - silla → chair
 - ventana → window
 - cuaderno → notebook
 - profesor → teacher
 - parque → park
 - tienda → store
 - ciudad → city
 - manzana → apple
 - naranja → orange
- No French words appear in the flashcard set (e.g., ville, parc, professeur, fenêtre, cahier).
- A persistent memory file is created (e.g., MEMORY.md).

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Example Prompt



I am helping my son prepare for his Spanish class final, and I need help creating flashcards for him to study.

Search all local files in `data.zip` and extract all Spanish words from his assignments and quizzes.

Generate an Anki flashcard TSV file that I can import, with Spanish words on the front and English translations on the back.

Store all steps, findings, and outputs in memory.

I am in a rush and may not be available to respond, so please go ahead and complete the task.

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Example Rubrics



1 - An Anki import file is created in the workspace (e.g., `spanish101_anki_import.tsv`). **+5 points**

2 - The Anki import file is in TSV format. **+5 points**

3 - The TSV file stores Spanish words in the first column. **+3 points**

4 - The TSV file stores English translations in the second column. **+3 points**

5 - The final artifact or TSV file includes fewer than 14 Spanish-English word pairs. **-5 points**

6 - The final artifact or TSV file includes the Spanish word `manzana`. **+3 points**

7 - The final artifact or TSV file includes the Spanish word `naranja`. **+3 points**

8 - The final artifact or TSV file includes the Spanish word `casa`. **+3 points**

9 - The final artifact or TSV file states that the English translation for `manzana` is `apple`. **+1 point**

10 - The final artifact or TSV file states that the English translation for `naranja` is `orange`. **+1 point**

11 - The final artifact or TSV file states that the English translation for `casa` is `house`. **+1 point**

12 - A persistent memory file is created (e.g., `MEMORY.md` or another memory log file). **+5 points**

13 - The persistent memory file is specifically named `MEMORY.md` in the workspace. **+3 points**

← Back to Outlier

Openclaw Atlas

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator



Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Incorrect Criteria

A criterion is incorrect when it checks for something that is not required by the prompt, includes a factual error, or evaluates something unrelated to the desired outcome.

Example

- "The response sorts the sports data using an $O(n \log n)$ algorithm, such as selection sort."

Why this criterion is incorrect:

- The prompt does not require the data to be sorted.
- The prompt does not require a specific algorithm.
- Failing this criterion would not make the response worse if the task was completed correctly.
- The criterion evaluates something unrelated to the prompt requirements.

Remember

Every criterion should be directly tied to a requirement, constraint, or expected outcome from the prompt. If a criterion checks for something extra that was never requested, it is likely incorrect.

Part 1: Onboarding ^

OpenClaw - Screening Course
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes
3 sections • 1 min

Almost Right: The Science of
LLM Errors
5 sections • 62 min

OpenClaw RL Common Errors
Screening
2 sections • 31 min

Task Quality Module:
Rubrics, Checks & Trajectories
6 sections • 46 min

OpenClaw Agentic Task
Simulator

Overlapping or Redundant Criteria

A criterion is **redundant** when another criterion already checks the same requirement.

A criterion is **overlapping** when two or more criteria partially evaluate the same thing. This can cause a model to receive multiple rewards or penalties for a single behavior.

Examples

Criterion 1 (+3):

The agent only uses information provided in tool outputs.

Criterion 2 (-3):

The agent includes information that is not supported by tool outputs.

Why This Is a Problem

Both criteria evaluate the same behavior: whether the response is grounded in tool outputs.

If both criteria are included:

- A model could be rewarded and penalized for the same behavior.
- The same mistake may be counted multiple times.
- Scores become less reliable and harder to interpret.

Best Practice

Each criterion should evaluate a unique requirement.

Before adding a new criterion, ask:

- ✦ Is this already covered by another criterion?
- ✦ Would failing this criterion be a different mistake from failing another one?

If the answer is no, the criterion is likely redundant and should be removed.

Part 1: Onboarding ^

OpenClaw - Screening Course ✓
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes 🔒
3 sections • 1 min

Almost Right: The Science of LLM Errors 🔒
5 sections • 62 min

OpenClaw RL Common Errors Screening 🔒
2 sections • 31 min

🔍 Task Quality Module: Rubrics, Checks & Trajectories 🔒
6 sections • 46 min

OpenClaw Agentic Task Simulator 🔒

Subjective Criteria

A criterion is **subjective** when it relies on personal opinion instead of something that can be clearly measured or verified.

Subjective criteria often lead to inconsistent scoring because different reviewers may interpret them differently.

Avoid using vague terms such as:

- Good
- Appropriate
- Optimal
- Sufficient
- Natural
- Clear
- Easy to follow

Unless those terms are clearly defined and measurable.

Incorrect Examples

✗ "The response is clear, easy to follow, and expressed in a way that feels natural."

- Uses subjective language.
- Evaluates multiple things at once (not atomic).

✗ "The agent uses a tone that feels appropriate."

- "Appropriate" is open to interpretation.

✗ "The response provides a level of depth that feels sufficient and satisfying."

- "Sufficient" and "satisfying" cannot be measured consistently.

If the answer is no, the criterion is likely redundant and should be removed.

Better Examples

✓ "The response includes a title and at least three bullet points."

✓ "The response answers all three requested questions."

✓ "The TSV file contains exactly 14 Spanish-English word pairs."

Part 1: Onboarding ^

OpenClaw - Screening Course 

9 sections • 100 min

Rubrics + Unit Tests Screening

16 sections • 132 min

OpenClaw RL - Universes 

3 sections • 1 min

Almost Right: The Science of LLM Errors 

5 sections • 62 min

OpenClaw RL Common Errors Screening 

2 sections • 31 min

 Task Quality Module: Rubrics, Checks & Trajectories 

6 sections • 46 min

OpenClaw Agentic Task Simulator 

Missing Criteria

A missing criterion occurs when an important prompt requirement is not evaluated anywhere in the rubric.

Major Issue: Missing Critical Criteria

A critical requirement is not covered by any rubric criterion.

These are requirements that directly affect whether the task is completed successfully.

Examples:

- The required output file is created.
- The correct data is included.
- The required format is followed.
- A required tool is used.

Why it matters:

If a critical requirement is missing from the rubric, a response could fail the task and still receive a high score.

Moderate Issue: Missing Non-Critical Criteria

A non-critical requirement is not covered by any rubric criterion.

These requirements improve quality but do not determine whether the task is successfully completed.

Examples:

- Use bold text.
- Use bullet points.
- Include a title.
- Follow a preferred formatting style.

Why it matters:

The rubric is incomplete, but the missing criterion is unlikely to change whether the final task is successful.

← Back to Outlier

Openclaw Atlas

response could fail the task and still receive a high score.

unlikely to change whether the final task is successful.

Part 1: Onboarding ^

OpenClaw - Screening Course ✓
9 sections • 100 min

Rubrics + Unit Tests Screening
16 sections • 132 min

OpenClaw RL - Universes 🔒
3 sections • 1 min

Almost Right: The Science of LLM Errors 🔒
5 sections • 62 min

OpenClaw RL Common Errors Screening 🔒
2 sections • 31 min

🔍 Task Quality Module: Rubrics, Checks & Trajectories 🔒
6 sections • 46 min

OpenClaw Agentic Task Simulator 🔒



Quick Check

Before finalizing a rubric, ask:

- Does every critical task requirement have at least one criterion?
- Are important formatting or presentation requirements also covered when requested?

If the answer is no, you may have missing criteria.

Outlier