


Rubrics for Atlas



## By the end of this practice task you will:

**Define a Rubric** generally.

Explain how to **Apply a Rubric** to this Project specifically.

**Define a Criterion** in the context of this Project.

**Assign a Weight** to each Criterion based on impact.

**Avoid** common mistakes.

**Complete a Rubric** for an actual Task with live feedback!





### A rubric is simply a checklist of criteria attached to a prompt.

To compose a rubric, a data annotator starts with a **prompt**.

The annotator then imagines what an **ideal response** to that prompt would look like.

Finally, the annotator breaks down the **features** of that ideal response into a set of **binary (yes/no) criteria**.

For example, if the prompt was **"Who was the second president of the United States? Just the name."**

You could create a **simple Rubric** by writing two criteria:

**Criterion #1**

The response identifies the second president of the United States as John Adams.

Taken together, this Rubric defines the output. Consider responses from two different models:

✓ **1:** John Adams.



**Criterion #2**  
The response contains only a name.

✗

**2:** The second president of the United States was John Adams.

- Why is this output worse?

While both outputs are technically correct, this Rubric allows us to evaluate the model on all of the specific ways it can succeed or fail to meet the user's expectations.

The fundamental concept is fairly straightforward, but things get more interesting when we attempt to create a Rubric for deep reasoning or complex Output, like what we expect from Atlas.

Let's learn how ↓

Applying Rubrics to Atlas

For this project, we care about evaluating the **Output** of a **very specific request**.

This means a Rubric you generate for one Task will **not be applicable to any other Task**.

Each Rubric you create will Evaluate:

**Agent Architecture and Behavior**  
Did the model complete the core objective?

**Reasoning Quality**  
Did the model follow required formatting and decision rules?

**Final Artifact/Output Quality**  
Did the model produce the required artifact exactly as described?

**Outcome Based Evaluation:** Models will be **evaluated primarily by the results they produce**, it's critical that the desired **Output is something that can be clearly defined**, not vague or unverifiable. *Consider the difference between a request for "lunch" and a request for a "grilled cheddar cheese sandwich on sourdough bread".*

In order to evaluate a models Output correctly, we need high quality Criteria outlining each specific expectation we have of the model. We will start by defining "High Quality Criteria".

Let's go ↓

Defining Criteria

Ideal criteria can vary from greatly from project to project. For Atlas, individual criteria will meet the expectations below. Let's use this prompt as an example: **"Create a PDF meal itinerary for my trip to New York City from June 10-19, include lunches and dinners. I'll also need the airport code for LaGuardia so I can track my flight."**

**Self-Contained**  
You can determine if a single Criterion is met or not met based solely on the model response.  
**Ask:** Is external research necessary?

**Atomic**  
A Criterion should evaluate only one component of a response.  
**Ask:** Could the model get this criterion partially right?  
**Bad Example:** The Itinerary is a PDF for dates 6/10-6/19

**Objective**  
Vague or subjective qualifiers (e.g., "good," "appropriate," "complete") are used in the criterion.  
**Ask:** Is this a verifiable truth or just my opinion?

**Categorized**  
Every criterion must closely match one of the following:  
→ **Task Completion:** Is the task complete?  
→ **Instruction Following:** Was the entire ask satisfied?  
→ **Exhaustive/Unrestricted:**

**Bad Example:** The model states the airport code for LaGuardia.

**Good Example:**

The model states **LGA** as the airport code for LaGuardia.

**PDF** for a trip to NYC.

**Good Example:**

1. The Itinerary is in **PDF** format.
2. The Itinerary is for a **trip to NYC**.

**Bad Example:** The itinerary lists events in a logical format.

**Good Example:**

The itinerary lists events in **chronological order**.

**Factuality/Invention:** Does the output contain false or made up info?

→ **Tool Use:** Does the model use expected tools?

→ **Agent Behavior:** A broad, open category.

## An Exception to Atomicity: Unit Tests

Usually, **Atomicity** requires you to **check for individual instances of a requirement**, for example:

1. *Itinerary includes lunch on the first day of the trip.*
2. *Itinerary includes lunch on the second day of the trip.*
3. *Itinerary includes....*

This is great, until you find yourself **assessing extremely large numbers of repeated actions** (ex: adding a years worth of expenses to a spreadsheet or sending a marketing email to 1000 recipients).

For this Project, you will assess repeated actions with **Unit Tests** consisting of 4 Criteria:

**Aggregate Count Criterion:** *Annual Itinerary evaluation includes 283 lunch events.* - **Not Atomic, but ok here**

**Random Spot Check 1:** *Annual Itinerary evaluation includes lunch on 2/2/2025.* - **Atomic**

**Random Spot Check 2:** *Annual Itinerary evaluation includes lunch on 5/7/2025.* - **Atomic**

**Random Spot Check 3:** *Annual Itinerary evaluation includes lunch on 12/10/2025.* - **Atomic**

This is an exception

## Weighting Criteria

### 🔴 Negative Criteria

These are criteria that take points away from the model for **actively** doing something detrimental to the output, **you must include at least one**.

- 1: Slightly detrimental architectural or behavioral flaw.
- 3: Detrimental to reliability but doesn't invalidate output.
- 5: Critically detrimental to results or invalidates output.

**Bad Example: -5:** The output **does not only include** events for the trip to New York City in June.

**Good Example: -5:** The output **includes** events outside of the trip to New York City from June 10-19.

### 🟢 Positive Criteria

Concise and Focused, these criteria **cover every critical step needed to achieve the desired outcome**.

- +1: Slightly important polish to an otherwise correct outcome.
- +3: Important to core competence and execution but not to outcome validity.
- +5: Critically important to outcome validity.

**Bad Example: +1:** The model generates a PDF output.

**Good Example: +5:** The model generates a PDF output.

Now that we've covered the fundamentals, it's time to put them to use! The following task is designed to guide you through the creation of a perfect Atlas Rubric.

You will receive live feedback as you go.

Let's go ↓