

← Back to Outlier

Opportunity Plover

Part 1: Onboarding ^

Opportunity Plover Onboarding Course

15 sections • 70 min

Opportunity Plover Post-Workshop Quiz

3 sections • 21 min



Opportunity Plover Common Errors

18 sections • 1 min



Part 2: Tasking 🔒



Part 3: Rubric Generation

Part 3: Generate the Rubric 🚀

You must create an atomic and self-contained rubric:

- The rubric consists of multiple criteria, each evaluating a different aspect of the model's answer and trajectory.
- Each criterion must be simple, self-contained, and cover only one step separately.
- Criteria must be understandable without the prompt or external searches, using only the output of a single tool 🛠️.

This ensures the auditor can evaluate the model fairly, based only on the rubric and one tool's output

Once you have the rubric, you must evaluate the SOTA model's answer using all the criteria: **For each criterion, check if the model fulfills it or not.**

Part 1: Onboarding ^

Opportunity Plover Onboarding Course

15 sections • 70 min

Opportunity Plover Post-Workshop Quiz

3 sections • 21 min

🔒

Opportunity Plover Common Errors

18 sections • 1 min

🔒

Part 2: Tasking 🔒 v

✔ Key qualities of each criterion

Each criterion in a rubric should be:

1. Clear & Complete

Meaning:

- The statement is written as a full sentence and communicates exactly what is being evaluated.
- It provides enough context to make the grader’s job effortless.

Why It Matters:

Ambiguity in rubric wording leads to inconsistent grading.

Example:

- “Check correctness of the result.” (too vague)
- “The model’s response must provide the total revenue in USD for 2024.”

2. Atomic (Single Focus)

Meaning:

- The criterion covers one and only one behavior or requirement.

Why It Matters:

Duplicated criteria (“The model must fetch and compare data”) make it unclear which part failed

Part 1: Onboarding ^

Opportunity Plover Onboarding Course

15 sections • 70 min

Opportunity Plover Post-Workshop Quiz

3 sections • 21 min



Opportunity Plover Common Errors

18 sections • 1 min



Part 2: Tasking 🔒



2. Atomic (Single Focus)

Meaning:

- The criterion covers one and only one behavior or requirement.

Why It Matters:

Bundled rubrics (“The model must fetch and compare data”) make it unclear which part failed.

Tip:

If a criterion uses 'and' or 'or', check if it can be split into two.

3. Specific (Binary and Objective)

Meaning:

- It can be evaluated with a simple yes or no.
- Avoid subjective or relative phrasing like “several,” “enough,” “appropriately,” etc.

Why It Matters:

Specificity ensures consistent LLM auto-grading and human reviewer agreement.

Example:

- “Model should include relevant details.”
- “Model must list at least three relevant details from the dataset.”

Part 1: Onboarding ^

Opportunity Plover Onboarding Course

15 sections • 70 min

Opportunity Plover Post-Workshop Quiz

3 sections • 21 min



Opportunity Plover Common Errors

18 sections • 1 min



Part 2: Tasking 🔒 v

4. Self-Contained

Meaning:

- The grader can understand and check it without looking at the prompt, tool outputs, or any external source.
- Every number, name, or reference required for evaluation appears in the criterion itself.

Why It Matters:

Rubrics are evaluated in isolation during audits; missing context causes scoring errors.

Example:

- “Model must compute the total as mentioned in the prompt.”
- “Model must compute the total sales value as 1,250 USD.”

5. Accurate

Meaning:

- The information stated in the criterion is factually correct and matches the tool output or GTFA.

Why It Matters:

Inaccurate rubrics lead to wrong scoring and mis-training signals for the model.

Example:

- “Model must use addition for averaging.” (incorrect logic)
- “Model must compute the average using sum divided by count.”

Part 1: Onboarding ^

Opportunity Plover Onboarding Course
15 sections • 70 min

Opportunity Plover Post-Workshop Quiz
3 sections • 21 min



Opportunity Plover Common Errors
18 sections • 1 min



Part 2: Tasking v

6. Categorized

Meaning:

- Each rubric must be labeled under the correct category and target.
 - Category: Tool Selection, Query Construction, Multi-step Planning, Answer Correctness, etc.
 - Target: Process/Reasoning (trajectory) or Outcome (final answer).

Why It Matters:

Proper labeling ensures the automated grading system and reviewers interpret the rubric's intent correctly.

7. Grounded in Task (Optional Add-on)

Meaning:

- The criterion should directly stem from what the prompt and task expect.
- Avoid rubrics that check extra or irrelevant behaviors.

Example:

- "Model must justify its reasoning style."
- "Model must identify the top three sales regions based on total revenue."

Closing Line:

When every criterion meets these six (or seven) qualities, your rubric becomes objective, reproducible, and model-trainable — the gold standard in Opportunity Plover quality control.

Part 1: Onboarding ^

Opportunity Plover Onboarding Course

15 sections • 70 min

Opportunity Plover Post-Workshop Quiz

3 sections • 21 min

Opportunity Plover Common Errors

18 sections • 1 min

Part 2: Tasking 🔒



Criteria Categories

When creating a rubric, it's important to understand the different dimensions you can evaluate. Each type of criterion focuses on a specific aspect of the model's performance. Rubrics are divided into two overarching targets:

- Process/Reasoning (Trajectory): How the model works — the tool choice, order, and reasoning steps.
- Outcome (Final Answer): What the model produces — the final results, correctness, and completeness.

Tool Selection (Process/Reasoning)

What it tests: Whether the model selects the correct tools required to solve the task and avoids redundant or invalid ones.

Common mistake: Including tools that don't align with the task or missing mandatory ones.

? Query Construction (Process/Reasoning)

What it tests: The model's ability to construct syntactically valid, schema-compliant API calls with correct parameter values.

Common mistake: Hardcoding or mislabeling parameters, or specifying only part of the required schema.

Multi-Step Planning (Process/Reasoning)

What it tests: Whether the model maintains logical order between tool calls, reuses earlier outputs correctly, and achieves intermediate goals.

Common mistake: Skipping dependencies or re-calling a tool unnecessarily instead of using existing results.

Part 1: Onboarding ^

Opportunity Plover Onboarding Course
15 sections • 70 min

Opportunity Plover Post-Workshop Quiz
3 sections • 21 min

Opportunity Plover Common Errors
18 sections • 1 min

Part 2: Tasking 🔒 v

✔ Answer Correctness (Outcome)

What it tests: Whether the model's final response matches the ground truth within defined tolerance (e.g., numeric $\pm 10\%$).

Common mistake: Stating partially correct information or adding fabricated details not grounded in tool outputs.

📌 Task Completeness (Outcome)

What it tests: Whether all required subtasks were executed — nothing skipped, omitted, or left unresolved.

Common mistake: Missing a required output or stopping after partial fulfillment of the prompt.

✎ Output Formatting (Outcome)

What it tests: Whether the final response follows proper structure — correct format, clear labeling, citations, and user readability.

Common mistake: Disorganized structure, missing fields, or formatting inconsistent with user expectations.

⊘ Negated Criteria (Advanced, Optional)

What it tests: Ensures the model avoids specific incorrect actions (e.g., calling irrelevant tools or introducing hallucinations).

Common mistake: Adding trivial or meaningless negations (e.g., 'Model must not use check_stock_prices' in a weather prompt).

Summary

Each rubric criterion should map clearly to one of these categories. Process rubrics measure how the model works. Outcome rubrics measure what the model produces. Together, they give a complete, balanced evaluation of model performance.

Part 1: Onboarding ^

Opportunity Plover Onboarding Course
15 sections • 70 min

Opportunity Plover Post-Workshop Quiz
3 sections • 21 min

Opportunity Plover Common Errors
18 sections • 1 min

Part 2: Tasking 🗝️ v

📌 Criterion Types

- **Essential** ✅
Must be met for the task to reach the Ground Truth Final Answer (GTFA).
If missing, the response fails.
Example: "The model must use the ``finance_get_stock_price`` tool to retrieve AAPL stock price."
- **Optional (Nice to Have)** ✨
Improves quality or clarity, but is not required to reach the GTFA.
Example: "The model explains why AAPL's stock is higher than TSLA's."
- **Negated** ❌
Describes actions the model must avoid; doing them makes the response incorrect.
Example: "The model must not call the ``weather_forecast`` tool in a finance task."

🎯 Target

Each rubric must also specify its **target** — what part of the model's behavior it is evaluating:

- **Process / Reasoning** 🧠
The rubric targets the model's **internal reasoning steps** or its **tool calls**.
(Focus: *how* the model works through the task.)
The format for this category should be "The model must..." or "The model must not..."
- **Outcome** 📝
The rubric targets the model's **final response** as delivered to the user.
(Focus: *what* the model ultimately produces.)
The format for this category should be "The model's response must..." or "The model's response must not..."

Part 1: Onboarding ^

Opportunity Plover Onboarding Course

15 sections • 70 min

Opportunity Plover Post-Workshop Quiz

3 sections • 21 min

🔒

Opportunity Plover Common Errors

18 sections • 1 min

🔒

Part 2: Tasking 🔒 v

Rubric: Common Errors and Best Practices

These are the most common errors we've found based on quality reads when reviewing the rubrics. We've included a brief explanation of each error and the best way to avoid them or improve the criteria.

Rubric: Criteria Not Atomic

This error involves **bundling too much information into a single check**.

- The criterion aims to evaluate multiple, distinct steps, parameters, or pieces of information simultaneously (e.g., checking three different tool parameters within a single criterion). **This makes it impossible to know which part of the criterion failed.**

✓ Best Practice:

- Each criterion **must evaluate one single, logical unit**. If the model must identify five different pieces of information, you should ideally have five separate criteria, one for each.

Example	Prompt
✗ Not Atomic:	"C16: The model's trajectory must identify the contact list of individuals who visited the pint shop more than twice as Evelyn Fothergill, Kaitlyn Nakamura, Akua Owusu-Agyemang, Reginald Fothergill, and Zara Saeed."
Ⓢ Atomic:	"C16: The model's trajectory must identify Evelyn Fothergill as a contact who visited the pint shop more than twice." "C17: The model's trajectory must identify Kaitlyn Nakamura as a contact who visited the pint shop more than twice." "C18: The model's trajectory must identify Akua Owusu-Agyemang as a contact who visited the pint shop more than twice." "C19: The model's trajectory must identify Reginald Fothergill as a contact who visited the pint shop more than twice." "C20: The model's trajectory must identify Zara Saeed as a contact who visited the pint shop more than twice."

Part 1: Onboarding ^

Opportunity Plover Onboarding Course
15 sections • 70 min

Opportunity Plover Post-Workshop Quiz
3 sections • 21 min

Opportunity Plover Common Errors
18 sections • 1 min

Part 2: Tasking 🔒



C20: The model's trajectory must identify Zara Saeed as a contact who visited the print shop more than twice.

Rubric: Criteria Not Self-Contained (Too Vague)

This is when the **criterion is too vague to be evaluated on its own.**

- The criterion is written vaguely, using placeholders or context-dependent phrases like "the correct value" or "the aforementioned items." It forces the reviewer to look at the tool outputs or other criteria to understand what's being checked.

✓ Best Practice:

- A reviewer must be able to evaluate the criterion just by reading it and the model's step. **It must state the exact expected value, string, or action that should be taken.**

Example	Prompt
✗ Vague:	"C6: The model must run a python script designed to analyze the 'Pet Care 2023 Weekly Financials.csv' file."
🎯 Self-Contained:	"C6: The model must run a python script that calculates and prints the sum of the 'Profit' column from 'Pet Care 2023 Weekly Financials.csv'."
✗ Vague:	"C24: The model must identify the contents of the notes."
🎯 Self-Contained:	"C24: The model must identify the string ' Quantum computing research notes ' from the output of the <i>read_file</i> tool."

Part 1: Onboarding ^

Opportunity Plover Onboarding Course

15 sections • 70 min

Opportunity Plover Post-Workshop Quiz

3 sections • 21 min

🔒

Opportunity Plover Common Errors

18 sections • 1 min

🔒

Part 2: Tasking 🔒 v

Rubric: Overly Specific Criteria (Too Strict)

This is the opposite problem of not being self-contained; **the criterion is too rigid.**

- The criterion **locks the model into one specific tool or one exact file name** when other valid tools or names could accomplish the same goal. This unfairly penalizes correct, alternative solutions.

✓ Best Practice:

- Write criteria that **evaluate the intent or outcome, not just one specific method**. If multiple tools can achieve the same result, the criterion should allow for both.

Example	Prompt
✗ Too Strict:	"C1: The model must use the <i>email_search_emails</i> tool." (Fails if <i>email_list_emails</i> could also work.)
🎯 Flexible:	"C1: The model must use a tool (like <i>email_search_emails</i> or <i>email_list_emails</i>) to find emails from Rohan."
✗ Too Strict:	"C5: The model must name the python file <i>analyze_pet_care.py</i> ." (The file name is irrelevant to the solution.)
🎯 Flexible:	"C5: The model must execute a python script that calculates the week with the lowest profit from the 'Pet Care 2023 Weekly Financials.csv' file."

Part 1: Onboarding ^

Opportunity Plover Onboarding Course
15 sections • 70 min

Opportunity Plover Post-Workshop Quiz
3 sections • 21 min

Opportunity Plover Common Errors
18 sections • 1 min

Part 2: Tasking 🔒

Lacking Distinct Coverage

- This fail state means **the prompt requires multiple, complex steps** (e.g., 4+ tool calls), **but the rubric has fewer than 10 distinct criteria to evaluate the model's performance**. You can't just have 10 criteria; they must each check a unique, logical step.
- For a complex task, the model must make many small, correct decisions (e.g., "list files," "identify the correct file," "read the file," "extract the value," "calculate a new value," etc.). **A rubric with too few criteria fails to capture and evaluate this complex reasoning process.**

✓ Best Practice:

Ensure any prompt that requires 4 or more tool calls has at least 10 distinct, non-redundant criteria.

- **What is "distinct"?** A criterion that checks a genuinely different logical step or action.
- **What is NOT "distinct"?** Having three criteria that check three different parameters for the exact same tool call (*C1: check param 'x', C2: check param 'y', C3: check param 'z'*) may only be counted as one distinct step (evaluating that single tool call). Similarly, using the same tool three times in a row to do the same action (like *filesystem_list_directory* to go 3 folders deep) is considered one logical "find the file" action, not three distinct ones.

Part 1: Onboarding ^

Opportunity Plover Onboarding Course
15 sections • 70 min

Opportunity Plover Post-Workshop Quiz 
3 sections • 21 min

Opportunity Plover Common Errors 
18 sections • 1 min

Part 2: Tasking v

Example	Prompt
✖ Lacking Coverage:	Prompt: "Find the config.py file in the /app/src/utils/ directory, read it, and tell me the value of the API_KEY variable." Rubric has only 7 criteria: 1. Tool: filesystem_list_directory (path='/app') 2. Tool: filesystem_list_directory (path='/app/src') 3. Tool: filesystem_list_directory (path='/app/src/utils') 4. Tool: filesystem_read_file (path='/app/src/utils/config.py') 5. Multi-step: Identify the line with API_KEY. 6. Tool: code_interpreter (to extract the value). 7. Outcome: State the correct API key.

Part 1: Onboarding ^

Opportunity Plover Onboarding Course

15 sections • 70 min

Opportunity Plover Post-Workshop Quiz

3 sections • 21 min

🔒

Opportunity Plover Common Errors

18 sections • 1 min

🔒

Part 2: Tasking 🔒 v

Example	Prompt
 Sufficient Coverage:	Prompt: "Find the config.py file in the /app/src/utils/ directory, read it, and tell me the value of the API_KEY variable." Rubric has 11 distinct criteria: 1. Tool: filesystem_list_directory (to start search). 2. Query: Path parameter is set to /app. 3. Multi-step: Identify src from the output. 4. Query: Path parameter is set to /app/src. 5. Multi-step: Identify utils from the output. 6. Query: Path parameter is set to /app/src/utils. 7. Multi-step: Identify config.py from the output. 8. Tool: filesystem_read_file. 9. Query: Path parameter is set to /app/src/utils/config.py. 10. Multi-step: Identify the string API_KEY = "xyz123" from the file content. 11. Outcome: State that the API key is "xyz123".

Part 1: Onboarding ^

Opportunity Plover Onboarding
Course
15 sections • 70 min

Opportunity Plover Post-
Workshop Quiz
3 sections • 21 min



Opportunity Plover Common
Errors
18 sections • 1 min



Part 2: Tasking 🔒



Errors related to the Golden Trajectory

Errors related to the Golden Trajectory

1. Fundamental Misinterpretation of the Prompt

✗ Error:

- This is when **the contributor fundamentally misunderstands the goal, and the entire trajectory is built on that flawed assumption.**
- The solution path is incorrect from the start because the contributor misinterpreted a key requirement, treating a detail as a filter or vice versa.

✓ Example

- The contributor built the entire trajectory around filtering for "kittens allowed." However, the prompt only asked to "know whether they accept my kitten," not to exclude apartments that don't. This misalignment "skewed" the entire rubric.

Part 1: Onboarding ^

Opportunity Plover Onboarding
Course
15 sections • 70 min

Opportunity Plover Post-
Workshop Quiz
3 sections • 21 min



Opportunity Plover Common
Errors
18 sections • 1 min



Part 2: Tasking 🔒



Errors related to the Golden Trajectory

2. Incomplete or Factually Incorrect Ground Truth

✖ Error:

- This is when **the contributor's trajectory and final answer (GTFA) simply contain the wrong information or are missing required data.**
- The trajectory includes factually incorrect data (like wrong names, numbers, or IDs) or fails to find all the information the prompt asked for.

✔ Example

- The trajectory was missing critical data. The contributor did not account for the meetings with ids: [...], or the contributor's GTFA does not include the two additional apartments requested. **This means the "golden" answer itself was wrong.**

← Back to Outlier

Opportunity Plover

Part 1: Onboarding ^

Opportunity Plover Onboarding
Course
15 sections • 70 min

Opportunity Plover Post-
Workshop Quiz
3 sections • 21 min



Opportunity Plover Common
Errors
18 sections • 1 min



Part 2: Tasking 🔒



Errors related to the Golden Trajectory

3. Logically Flawed or Irrelevant Steps

✗ Error:

- This is when **the trajectory includes steps that are either unnecessary or make no logical sense in the context of the prompt.**
- The contributor's solution path includes tool calls or actions that do not help answer the request or are even counterproductive.

✓ Example

- The first 6 criteria are completely irrelevant to the task because they checked for a job title when the model needed an email address.

Opportunity Plover

Prompt Evaluation

Evaluate each of the provided prompts/GTFA pairs below and note down any errors.

Remember that prompts written for this project need to adhere to the following guidelines:

- Should have a single, ground-truth final answer (GTFA) that remains constant over time. Please note that data from all tools (except web search) is static
- Should require tool calls, i.e., cannot be answered through basic reasoning or general knowledge
- Should have no explicit mention of tool names (e.g., get_weather)
- Should sound natural, i.e., you should be able to imagine someone in real life asking it
- Should be clear and contain all relevant details **without conflicting instructions**
- Should be feasible given the available tool set

Remember to be as explicit as possible and state the reason why the prompt does not meet certain characteristics